

# Graph Neural Networks

**Shah Rukh Qasim**

Physik Institute, University of Zurich

[shah.rukh.qasim@cern.ch](mailto:shah.rukh.qasim@cern.ch)

29.10.2024

Maurizio Pierini (CERN)


Thea Aarrestad (ETH Zürich)

# What are we doing today?

- Graph Neural Networks (GNNs)
- Part 1 (Lecture)
  - Why GNNs?
  - And then quickly go to a simple modern GNN / with message passing
    - So you can already get the understanding to be able to use GNNs
  - And then we'll have a little bit of a history lesson
  - And possibly discuss some advanced topics as well
- Part 2 (Tutorial)
  - We'll train a GNN for molecule classification on collab


# Why GNNs?

- A lot of data cannot be represented as fixed structure
  - That is sequences
  - or grids
- So we need a more generic data structure: graphs
- Graphs:  $G = (V, E)$ 
  - Nodes / vertices ( $V$ ): relate to an “object” in your data
  - Edges ( $E$ ): captures the relationship between
  - Let's say node features are  $h_u$  (corresponding to the node  $u$ )
  - And the edge features are  $e_{vw}$  (corresponding to edge  $e$ )
- Graphs are very generic
  - Can represent sequences and images
  - And sets





Networks


VS.



Images




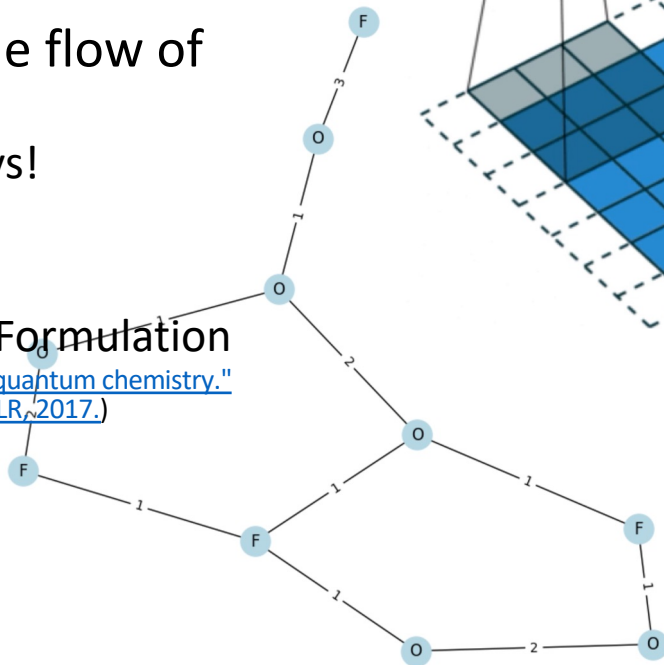
Mitchell, J. Clyde. "Social networks." Annual review of anthropology 3 (1974): 279-299.



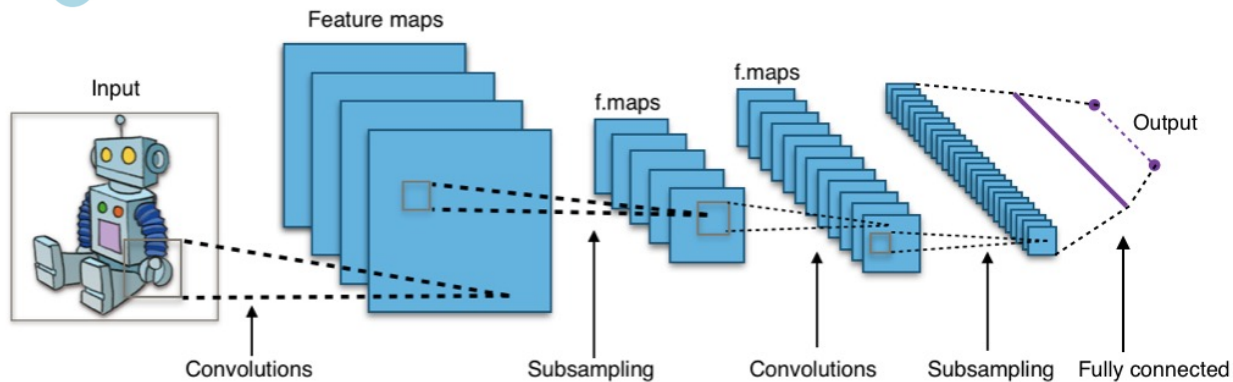

Text

# GNNs

- GNNs: Pairing up the graph structure with neural networks
- Use the graph structure for the flow of information
  - Can also be used in other ways!
- In the next slides:
  - We'll adapt Message Passing Formulation  
([Gilmer, Justin, et al. "Neural message passing for quantum chemistry." International conference on machine learning. PMLR, 2017.](#))




# GNNs and CNNs




# Graph Message Passing: Message Function

- Ignore edge features at the moment
- Every node has a feature vector  $h_n^t$
- Need to transform it into a message to send to its neighbours
- How?
- Apply a dense layer to it
- This is called a message function ( $M_t$ )
  - Message =  $M_t(h_n^t)$
- Are the weights different for different nodes?
  - No



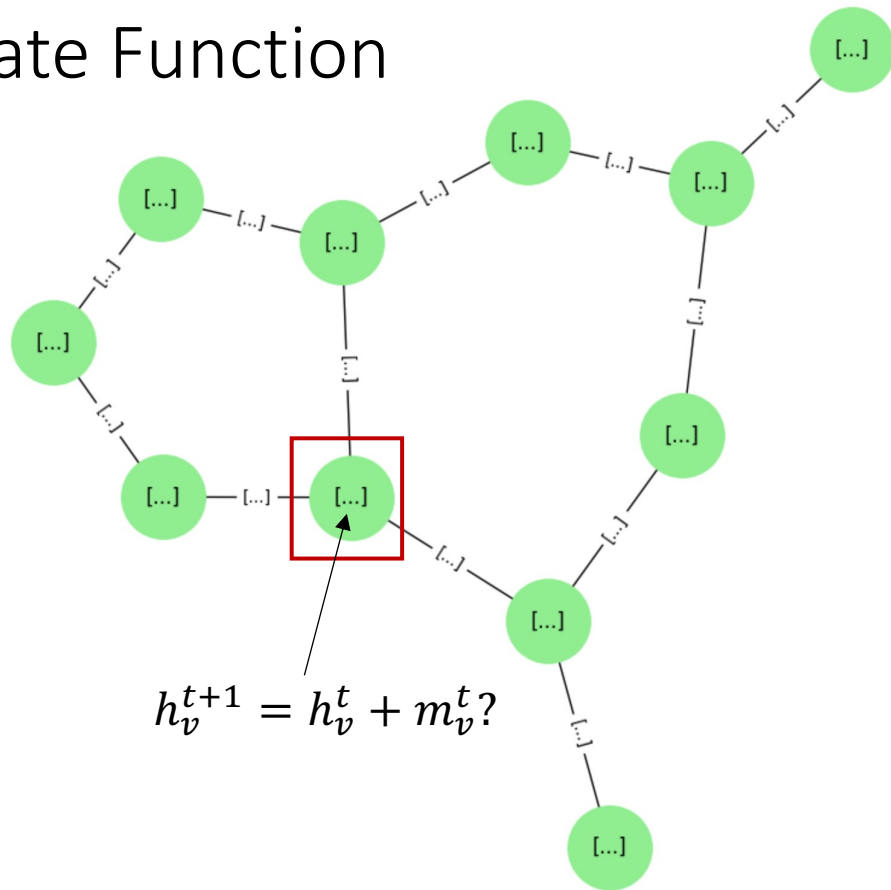
# Graph Message Passing: Aggregation function

- Every node ( $v$ ) collects messages from its neighbours ( $w$ )
  - $m_v^t = \prod_{w \in \mathcal{N}(v)} M_t(h_w^t)$
  - $\prod$  is not multiplication
- What is  $\prod$  ?
  - This is called **aggregation function**
  - We need it to be permutation invariant
  - Easiest and most common: sum!
  - $m_v^t = \sum_{w \in \mathcal{N}(v)} M_t(h_w^t)$
- Other aggregation functions:
  - Mean, min, max



# Graph Message Passing: Update Function

- Let's look at node  $v$  that has collected messages from its neighbours as  $m_v^t$
- The nodes previous features were:  $h_v^t$
- How do we merge them together for the next layer ( $t + 1$ )?
  - $h_v^{t+1} = U_t(h_v^t, m_v^t)$
- What is  $U_t$ ?
  - Again, a dense layer!
  - It is called an update function
- Before you apply a dense layer, how do you join the two vectors together ( $h_v^t, m_v^t$ )?
  - Concatenation
- You can also use only concatenation as the update function
  - But it will blow up the number of features as you add more layers (not ideal)





# Graph Message Passing

- Generalization: A message passing function ( $M_t$ ) operates on the hidden features of the source node, **the destination node and the features of the edges between them**
- The  $M_t$  is generally a dense layer
- Mean / min / max / sum are aggregation functions (permutation invariant)


$$m_v^{t+1} = \sum_{w \in N(v)} M_t(h_v^t, h_w^t, e_{vw})$$

- And then the vertex features are updated
  - simplest: also, a dense layer

$$h_v^{t+1} = U_t(h_v^t, m_v^{t+1})$$

- GRUs have been tried too

[Li, Yujia, Tarlow, Daniel, Brockschmidt, Marc, and Zemel, Richard. Gated graph sequence neural networks. ICLR, 2016.](#)



# GNN tasks

## ● Graph level


- Example: Molecular graph classification
- Pool information from all the nodes
  - Mean / Min / Max
- Output: The pooled vector
  - Also the loss function here
  - Generally after a couple of dense layers from the pooled information

## ● Node level


- Example: segmentation
- Output: the feature vector at the last layer of your GNN

## ● Edge level

- Example: In a social network, if two people are friends or not
- Don't pool information from the graph
- Get the output at every edge
  - If the edge features don't exist, do subtraction or concatenation




[Martínez, J.A., Cerri, O., Spiropulu, M., Vlimant, J.R. and Pierini, M., 2019. Pileup mitigation at the Large Hadron Collider with graph neural networks. The European Physical Journal Plus, 134\(7\), p.333.](#)



[Qasim, Shah Rukh, Hassan Mahmood, and Faisal Shafait. "Rethinking table recognition using graph neural networks." 2019 International Conference on Document Analysis and Recognition \(ICDAR\). IEEE, 2019.](#)

# Graph Neural Networks: Practical Considerations

- Use one or more dense layers without message passing at the start
- Use one or more dense layers at the end after message passing
- Don't have to use a single dense layer as message passing function
  - Can use more for a more complex function
  - Balance between the complexity and the frequency of the messages
- Can play with concatenation
- What stays the same?
  - Almost everything apart from the computation graph
    - Mini batch training
    - Train/validation split
    - Batch norm/dropout
    - Loss function
    - Adam optimizer
    - etc




An example of a GNN architecture

[Martínez, J.A., Cerri, O., Spiropulu, M., Vilmant, J.R. and Pierini, M., 2019. Pileup mitigation at the Large Hadron Collider with graph neural networks. The European Physical Journal Plus, 134\(7\), p.333.](#)


# History: Spectral Networks

- [Bruna, Joan, Wojciech Zaremba, Arthur Szlam, and Yann LeCun. "Spectral networks and locally connected networks on graphs." arXiv preprint arXiv:1312.6203 \(2013\).](#)
- They tried to generalise CNNs beyond the regular array dataset paradigm
- They replaced the translation-invariant kernel structure of CNNs with hierarchical clustering
- Introduced spectral convolutions
  - Done using the laplacian matrices of the graph
  - Later extended by [Kipf, T.N. and Welling, M., 2016. Semi-supervised classification with graph convolutional networks. arXiv preprint arXiv:1609.02907.](#)
- Today you won't use the spectral domain and perform message passing directly in the spatial domain



# History: Message Passing

- Traced back to [Duvenaud, David K., et al. "Convolutional networks on graphs for learning molecular fingerprints." Advances in neural information processing systems 28 \(2015\).](#)
- Introduced “convolutional neural network that operates directly on graphs”
- Different language, similar idea



---


## Algorithm 2 Neural graph fingerprints

---

- Input:** molecule, radius  $R$ , **hidden weights**  $H_1^1 \dots H_R^5$ , **output weights**  $W_1 \dots W_R$
  - Initialize:** fingerprint vector  $\mathbf{f} \leftarrow \mathbf{0}_S$
  - for** each atom  $a$  in molecule
  - $\mathbf{r}_a \leftarrow g(a)$       ▷ lookup atom features
  - for**  $L = 1$  to  $R$       ▷ for each layer
  - for** each atom  $a$  in molecule
  - $\mathbf{r}_1 \dots \mathbf{r}_N = \text{neighbors}(a)$
  - $\mathbf{v} \leftarrow \mathbf{r}_a + \sum_{i=1}^N \mathbf{r}_i$       ▷ **sum**
  - $\mathbf{r}_a \leftarrow \sigma(\mathbf{v} H_L^N)$       ▷ **smooth function**
  - $\mathbf{i} \leftarrow \text{softmax}(\mathbf{r}_a W_L)$       ▷ **sparsify**
  - $\mathbf{f} \leftarrow \mathbf{f} + \mathbf{i}$       ▷ **add to fingerprint**
  - Return:** **real-valued** vector  $\mathbf{f}$
-


# What about point clouds?

- Your non structured data is a bunch of points
  - aka a point cloud
- No graph structure defined
  - Very important for science, especially for physics
  - A data from a calorimeter or a tracker doesn't have a graph defined
  - A set of particles in a jet don't have a graph defined
- How do I apply machine learning here?




# GNNs on Point Clouds

- Answer 1:
  - Connect everything to everything
  - Introduce virtual nodes
  - Will not scale!
- Answer 2:
  - Build a graph locally in
    - Radially in the input domain...
    - Or k Nearest Neighbour (kNN)



Fully connected graph



kNN graph


# Better answer: Dynamic GNNs

- Build the graph in a dynamically learned space
- Feature vector for every point
- You apply a neural network to feature vectors of every point to get feature space  $F_s$
- Build a graph as KNN of every node using euclidean distance in  $F_s$ 
  - So you have  $N \times K$  edges
  - Perform message passing with permutation invariant aggregation functions
  - Dynamic GNN:
    - The KNN graph isn't built in the original space but in a learned feature space



# Dynamic GNNs: Edge Conv

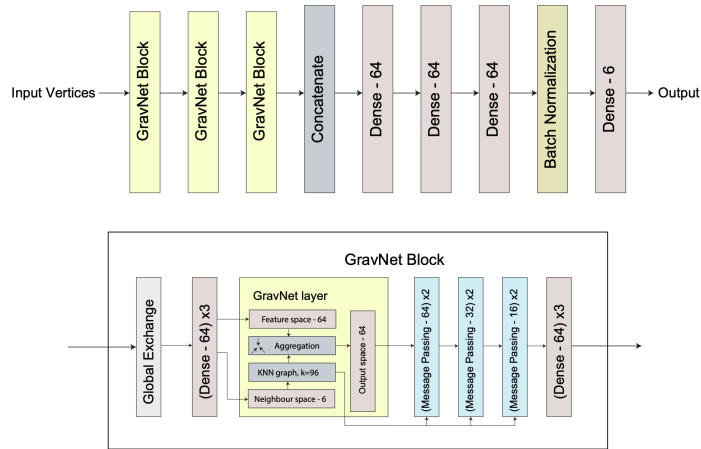

- kNN graph is built in  $h_t$  or  $x^t$
- And then the next set of hidden features are constructed as:
  - $x_i^{t+1} = m_i^{t+1} = \sum_{j \in \mathcal{N}(i)} M_t(x_i^t, x_j^t - x_i^t)$
- Can also use min / max functions
- Same story: Do it a bunch of times to build your GNN
  - Different: The graph is changed at every layer




Wang, Yue, et al. "Dynamic graph cnn for learning on point clouds." ACM Transactions on Graphics (tog) 38.5 (2019): 1-12.

# Dynamic GNNs: GravNet

- Edge Conv:
  - You build the graph in the feature space at step  $t$
  - Hope that the graph is good
    - Works very well in practice
- Can we do better?
  - At every layer, transform the features into two different spaces
    - Coordinate space ( $S$ ) – Low dimensional
    - Feature space ( $F_{LR}$ ) – Higher dimensional
  - $F_{LR}$  is weighed by (neg exp. of) distance between nodes in  $S$
  - In this way, the coordinate space is learned such that the nodes are brought closer if it helps the gradient
- Also much more performant:
  - kNN building is very expensive: dominates the time taken
  - Much faster in smaller dimensions



Qasim, Shah Rukh, et al. "Learning representations of irregular particle-detector geometry with distance-weighted graph networks." The European Physical Journal C 79.7 (2019): 1-11.




# GNNs on large graphs


- GraphSAGE (SAmple and AGgregate)

[Hamilton, Will, Zhitao Ying, and Jure Leskovec. "Inductive representation learning on large graphs." Advances in neural information processing systems 30 \(2017\).](#)


- Sample in the neighbourhood
- And then do the aggregation
- Here, the authors also explored LSTMs
  - Not permutation invariant though
- Applied on large graphs: citation and reddit graphs



1. Sample neighborhood



2. Aggregate feature information from neighbors



3. Predict graph context and label using aggregated information

# Graph Attention Networks



- [Veličković, P., Cucurull, G., Casanova, A., Romero, A., Lio, P. and Bengio, Y., 2017. Graph attention networks. \*arXiv preprint arXiv:1710.10903\*.](#)

$$e_{ij} = a(\mathbf{W}\vec{h}_i, \mathbf{W}\vec{h}_j)$$

$$\alpha_{ij} = \text{softmax}_j(e_{ij}) = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}_i} \exp(e_{ik})}$$

$$\vec{h}'_i = \sigma \left( \sum_{j \in \mathcal{N}_i} \alpha_{ij} \mathbf{W}\vec{h}_j \right)$$

- "Attention" on messages
- Maurizio will go into attention in more detail on November 19 when discussing transformers
- The paper is more for reference here as an important GNN paper



# Summary

- GNNs are powerful and work on very different types of datasets
- And now also very popular
- Transformers overshadow everything
  - But in a lot of cases, GNNs make more sense
- Have been applied to a large sets of problems across multiple domains
  - Traffic forecasting
  - Molecule property prediction
  - Recommendation systems
  - Reconstruction in particle physics
  - Simulation in particle physics
  - the list goes on...

# GNN Hands-on Tutorial

# Software Libraries for the use of GNNs

1. PyTorch: PyTorchGeometric / PyG: <https://pyg.org/>
  - Academia standard
2. **TensorFlow: TF-GNN**
  - We are going to use this to keep compatibility
  - Can use tools like keras here
3. Deep Graph Library (DGL): Framework agnostic

# Copy the notebook

- [https://colab.research.google.com/drive/1rmeb4qllbZ6WHxeqBmVzahPwKyur7skR?usp=share\\_link](https://colab.research.google.com/drive/1rmeb4qllbZ6WHxeqBmVzahPwKyur7skR?usp=share_link)
- Shorter link:
  - [bit.ly/3YCFCTf](https://bit.ly/3YCFCTf)
- Copy to save
- For slides: go to the very last of the notebook where there is a link to the slides





# Quick primer to NetworkX

- Stop at **Checkpoint #1**
- NetworkX is a Python library for creating and analyzing graphs and networks
- Also provides tools for visualization with matplotlib
- Has a bunch of classical graph algorithms (matching, clustering etc)
- We are going to use it to plot our graphs

```
my_graph = train_graphs[5]
my_graph_nx = graph_tensor_to_networkx(my_graph)
nx.draw(my_graph_nx)
```

# Quick primer to NetworkX

- Stop at **Checkpoint #2**
- We plot with the function we provided
  - with ids and without ids
- Add another node
  - And connect it to one of the existing ones
- Plot it again

```
my_graph = train_graphs[5]
my_graph_nx = graph_tensor_to_networkx(my_graph)
# nx.draw(my_graph_nx)
plot_networkx_graph(my_graph_nx, plot_with_ids=False)
```

# Graph Spec

1. Stop at CHECKPOINT #3
2. Not always needed
  - a. keras needs it
3. Looks scary piece of code but it is not doing much
4. Simplified for you below
5. If you understand the basic concept
  - a. Ask chatgpt to figure out the technical details for you

# To start

1. Stop at **CHECKPOINT #4**
2. Define the input
3. The graph is then merged across batch elements
  - a. Nothing is happening here in reality

```
# Model building with Keras's Functional API starts with an input object
# (a placeholder for the eventual inputs). Here is how it works for
# GraphTensors:
input_graph = tf.keras.layers.Input(type_spec=graph_tensor_spec)

# IMPORTANT: All TF-GNN modeling code assumes a GraphTensor of shape []
# in which the graphs of the input batch have been merged to components of
# one contiguously indexed graph. (There are no edges between components,
# so no information flows between them.)
graph = input_graph.merge_batch_to_components()
```

# Simple transformations

1. Apply a simple linear transformations to the initial features
  - Both for the nodes and edges
  - Similar to applying 1x1 conv
2. No message passing here

```
# Nodes and edges have one-hot encoded input features. Sending them through
# a Dense layer effectively does a lookup in a trainable embedding table.
def set_initial_node_state(node_set, *, node_set_name):
    # Since we only have one node set, we can ignore node_set_name.
    return tf.keras.layers.Dense(node_dim)(node_set[tfgnn.HIDDEN_STATE])

def set_initial_edge_state(edge_set, *, edge_set_name):
    return tf.keras.layers.Dense(edge_dim)(edge_set[tfgnn.HIDDEN_STATE])

graph = tfgnn.keras.layers.MapFeatures(
    node_sets_fn=set_initial_node_state, edge_sets_fn=set_initial_edge_state)(
    graph)
```

# Add another dense layer in between

- Before checkpoint 5, add two dense layers

# Message Passing

- From the documentation

- This layer can compute a convolution over an edge set by applying the passed-in `message_fn` for all edges on the concatenated inputs from some or all of: the edge itself, the sender node, and the receiver node, followed by pooling to the receiver node.

# Message Passing: Tasks

- Add three message passing layers
- In the first one, we are only using the sender node features
- In the second one, we are using the sender node features and the edge features
- In the third one, we are using the sender node features, the edge features and the receiver node features



# Graph Message Passing

- Generalization: A message passing function ( $M_t$ ) operates on the hidden features of the source node, **the destination node and the features of the edges between them**
- The  $M_t$  is generally a dense layer
- Mean / min / max / sum are aggregation functions (permutation invariant)


$$m_v^{t+1} = \sum_{w \in N(v)} M_t(h_v^t, h_w^t, e_{vw})$$

- And then the vertex features are updated
  - simplest: also, a dense layer

$$h_v^{t+1} = U_t(h_v^t, m_v^{t+1})$$

- GRUs have been tried too

[Li, Yujia, Tarlow, Daniel, Brockschmidt, Marc, and Zemel, Richard. Gated graph sequence neural networks. ICLR, 2016.](#)



# Final Pooling

- We are performing a graph level task
- We need to pool information from all the nodes
  - to the graph context
  - **Possible task: Do max pooling and concat before getting the output**
- Easy
  - Just do mean | min | max of the features of the nodes
  - Ignore the edges
  - And then you can add one or more dense layers if you want

```
readout_features = tfgnn.keras.layers.Pool(  
    tfgnn.CONTEXT, "mean", node_set_name="atoms")(graph)
```

```
# Put a linear classifier on top (not followed by dropout).  
logits = tf.keras.layers.Dense(1)(readout_features)
```

```
# Build a Keras Model for the transformation from input_graph to logits.  
return tf.keras.Model(inputs=[input_graph], outputs=[logits])
```

# The rest is the same

- Build the model
  - And the cost function
  - Use the same Adam optimize
- And then train
- And print some results
- GNNs are only doing the message passing part